

Distributed Systems Paradigms

José Orlando Pereira

HASLab / Departamento de Informática
Universidade do Minho



2021/2022

Motivation

- Handle a large number of clients and requests with a single server
- The “c10k problem” in 1999:
 - <http://www.kegel.com/c10k.html>
- Examples:
 - financial, games, ...
 - notifications in mobile apps
 - machine-to-machine (M2M)

Case study

- Simple chat server:
 - Forward all messages to all clients
- Consider:
 - Large number of clients
 - Slow connections



First threaded solution

- For each connection:
 - Handler thread
- When reading, write to all other connections
- Use buffering:
 - At user level (streams): To minimize system calls
 - In the kernel (socket): To cope with slow readers

Sockets in java.net

```
ServerSocket ss=new ServerSocket(12345);

while(true) {
    Socket s=ss.accept();

    InputStream is=s.getInputStream();
    OutputStream os=s.getOutputStream();

    // i/o

    s.close();
}
```

Buffers in java.net

```
ServerSocket ss=new ServerSocket(12345);

while(true) {
    Socket s=ss.accept();

    InputStream is=new BufferedInputStream(s.getInputStream());
    OutputStream os=new BufferedOutputStream(s.getOutputStream());

    // i/o

    os.flush();
    s.close();
}
```

Needed to
actually write!

Memory


- Memory: n connections x messages in transit ($\sim n^2$)
 - Caused by data copying in stacked abstractions
 - Serialization!
 - Overhead in allocation and garbage collection
- Solution: Store transient data in reusable shared buffers
 - Pointers/indexes into statically allocated data

Sockets in java.nio

```
ServerSocketChannel ss=ServerSocketChannel.open();  
ss.bind(new InetSocketAddress(12345));  
  
while(true) {  
    SocketChannel s=ss.accept();  
  
    // i/o  
  
    s.close();  
}
```



Buffers in java.nio

- Buffer = Array + Indexes:



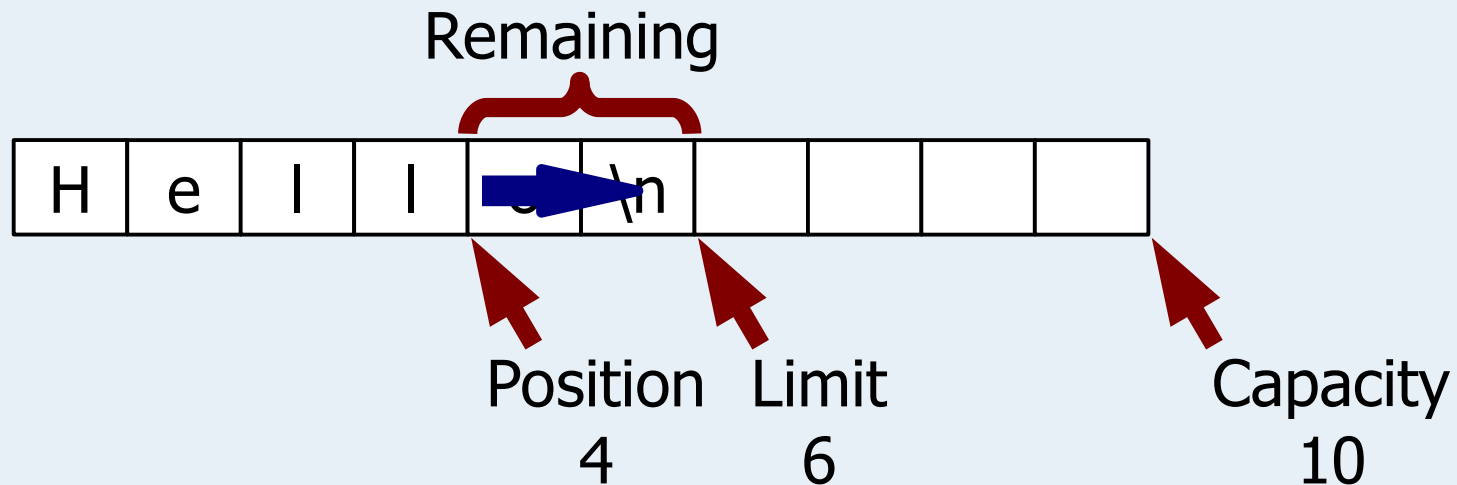
Buffers in java.nio

- Put/read: advances position, sets content




Buffers in java.nio

- Get/write: advances position, gets content

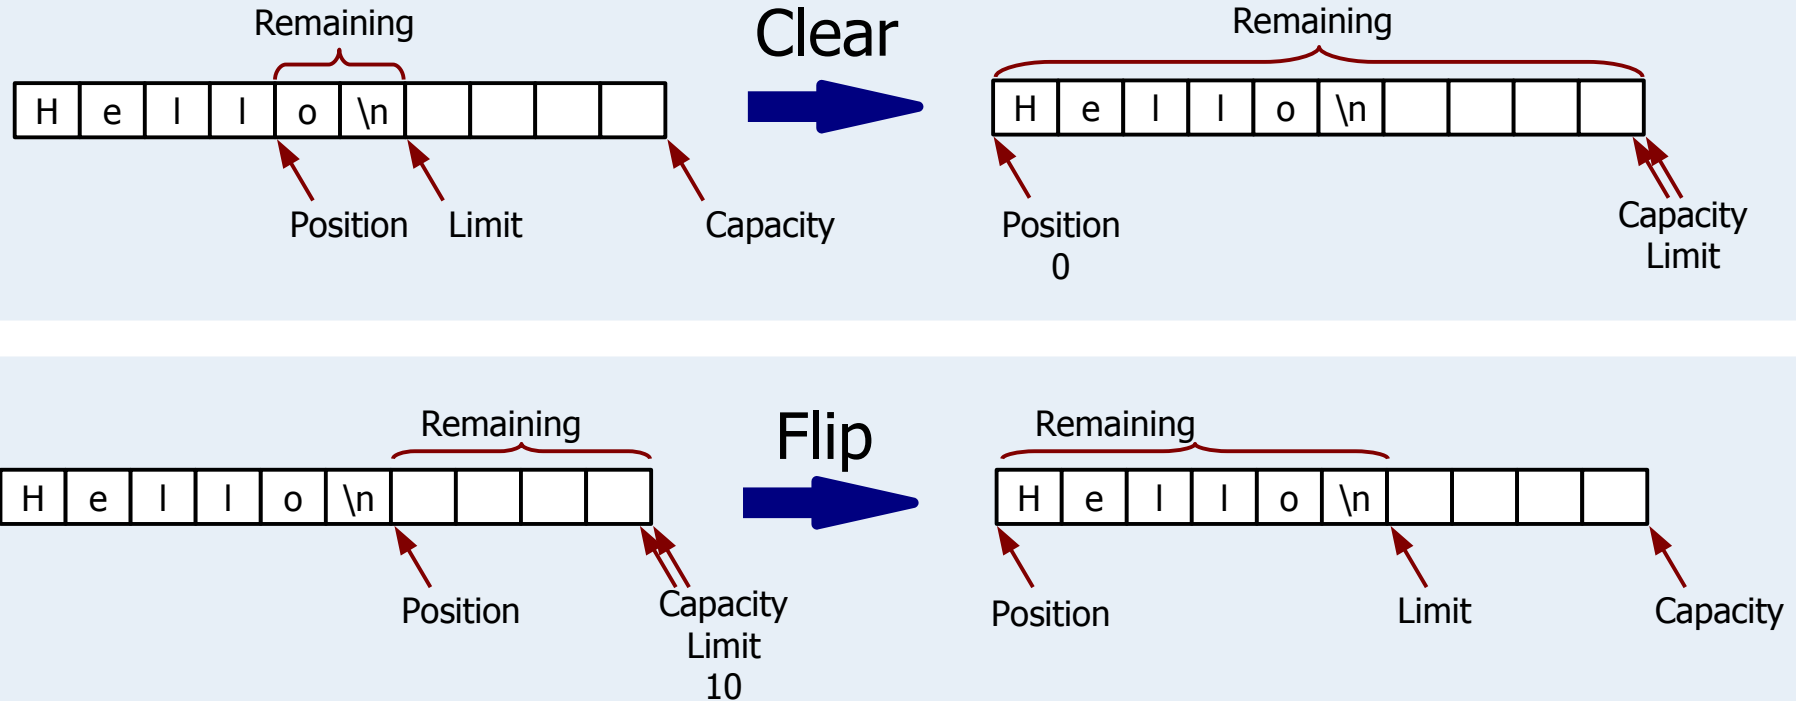


Buffers in java.nio



- Duplicate and wrap: multiple pointers into the same array



Buffers in java.nio



Buffers in java.nio



Sockets in java.nio

```
try {  
    ByteBuffer buf=ByteBuffer.allocate(100);  
  
    s.read(buf);  
    buf.flip();  
    for(SocketChannel r: receivers) {  
        r.write(buf.duplicate());  
    }  
} catch(IOException e) {  
    report(e);  
}
```

Shared buffers

- Memory used: messages in transit ($\sim n$)
- Ideally, never allocate or dispose of memory in normal operation:
 - No overhead, but...
 - Needs reference counting to know when to reuse

Flushing buffers

```
ByteBuffer buf=ByteBuffer.allocate(100);  
try {  
    s.read(buf);  
    buf.flip();  
    for(SocketChannel r: receivers) {  
        r.write(buf.duplicate());  
    }  
    buf.clear();  
} catch(IOException e) {  
    report(e);  
}
```

What if
write blocks?

Second threaded solution

- For each connection:
 - Reader thread + Pending queue + Writer thread
- When reading, insert in outgoing queues and notify writer threads
 - If a queue overflows, the reader must block, drop some data, or even disconnect the writer
- When writing, remove from queue
 - If readers might have blocked on this queue, notify them now

Threads summary

- Simple programming model
- Problems:
 - Memory overhead (stacks and buffers)
 - Context switches and lock contention
 - “Thundering herd”, hidden queue, and fairness