


# Event-driven I/O


- Instead of waiting for each I/O event with a dedicated thread...
- ...let the operating system explicitly notify the application of an I/O event

# Design issues

- Completion vs readiness



- Level-triggered vs edge-triggered



# Implementation issues

- Managing the event set
  - User vs. kernel-level
  - Examples: `select()/poll()` vs `epoll()`
- Control transfer
  - Blocking thread vs signals vs busy polling
  - Examples: `select()` vs `SIGIO` vs `DPDK/SPDK`

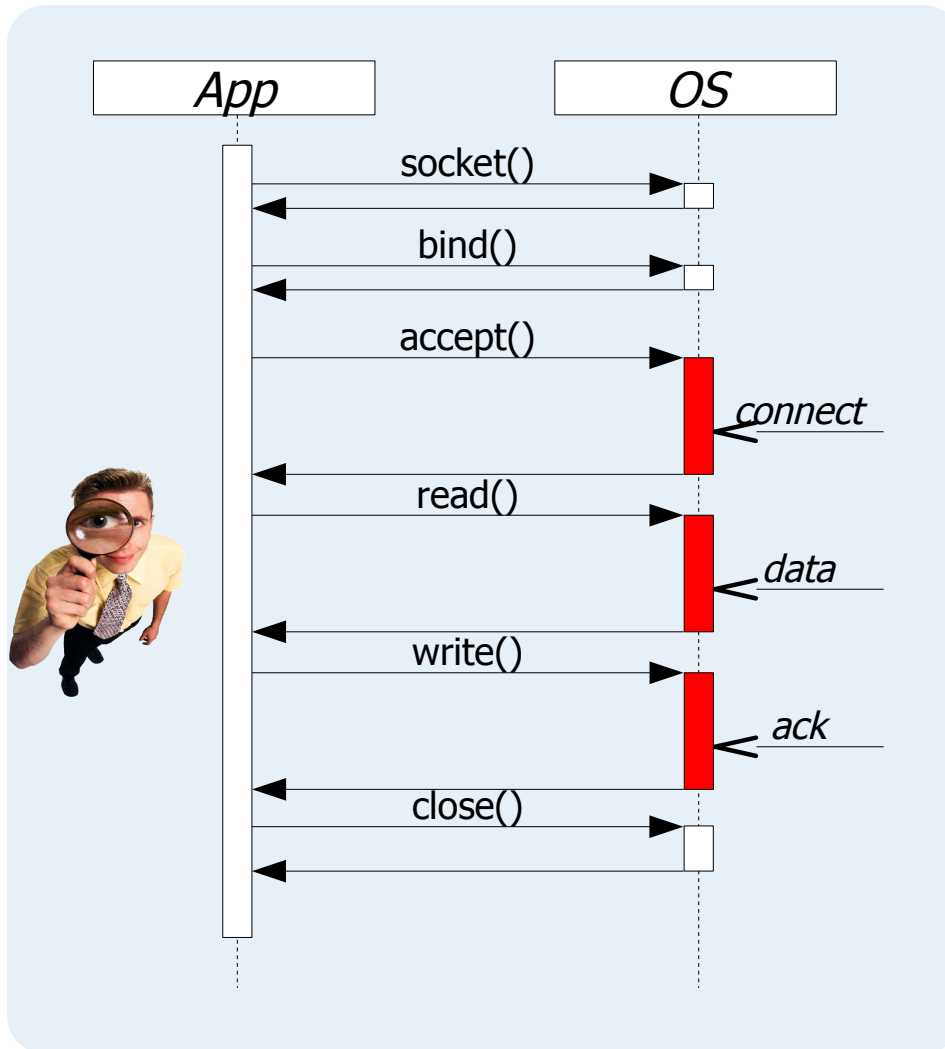
# Case studies in Java

- Asynchronous sockets (NIO2)
- Selectors (NIO)

# Asynchronous I/O

- For each blocking I/O operation, provide a callback to execute after the operation has completed
  - Completion event / edge-triggered
- General idea: Instead of:  
    `read(buf); doSomething();`  
do:  
    `read(buf, ()->{ doSomething(); })`

# Threaded version



```

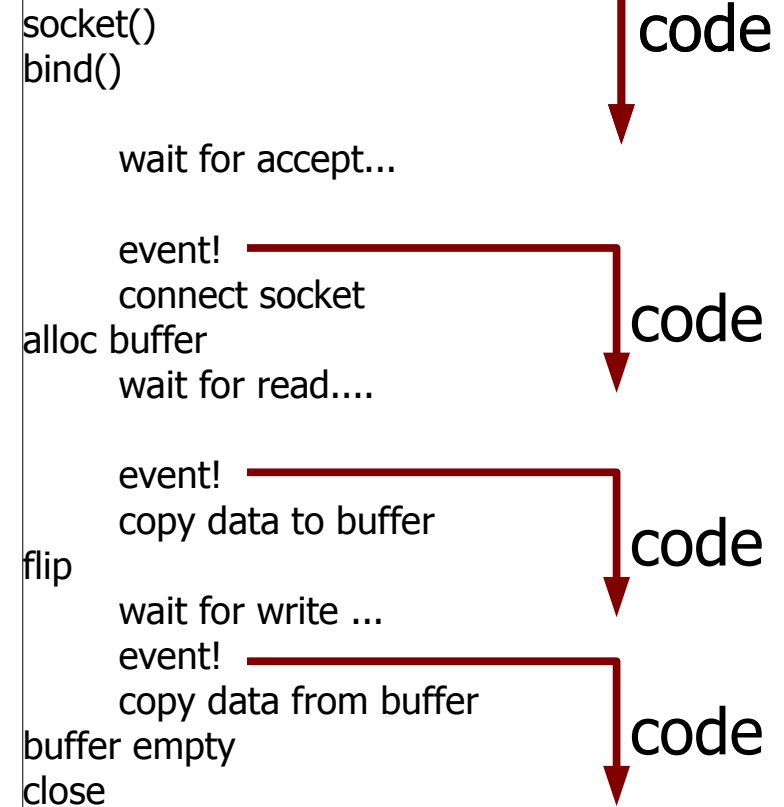
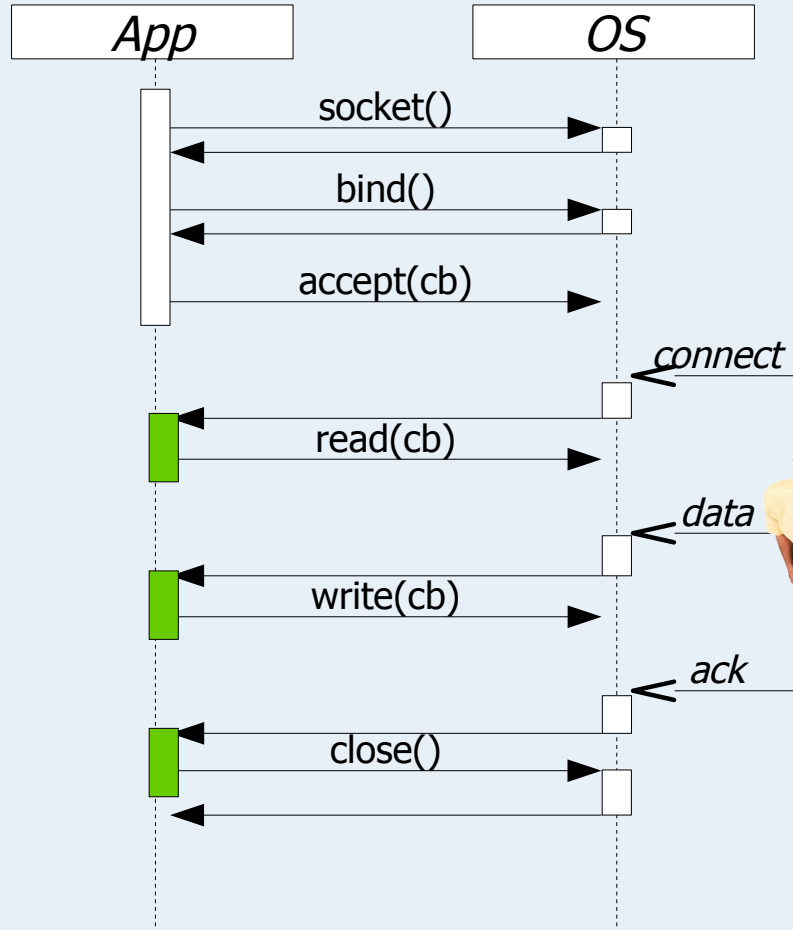
socket()
bind()
  wait for accept...

  event!
  connect socket
alloc buffer
  wait for read....
  event!
  copy data to buffer
flip
  wait for write ...
  event!
  copy data from buffer
buffer empty
close
  
```

code



# Asynchronous version



# Inversion of Control (IoC)

- With threads:
  - The program controls flow
  - Calls into the framework for specific tasks
- With events:
  - The framework controls flow
  - Calls back the program for specific tasks



# Asynchronous I/O

- Avoids having a dedicated thread for each event source
- However:
  - Requires captive memory for idle I/O channels
  - Hides threading policy within the framework
- Available in Java with NIO2 `AsynchronousSockets`

# Blocking sockets

```
try {  
    ByteBuffer buf=ByteBuffer.allocate(100);  
  
    s.read(buf);  
    buf.flip();  
  
    r.write(buf);  
} catch(IOException e) {  
    report(e);  
}
```

# Translation to CompletionHandler

```
try {  
    C c = codeBefore(...);  
  
    R r = operation(...);  
  
    codeAfter(c, r);  
} catch(Exception e) {  
    handleException(e);  
}
```

```
C c = codeBefore(...);  
asyncOperation(..., c, new CompletionHandler<R,C>() {  
    public void sucess(R r, C c) {  
        codeAfter(c, r);  
    }  
    public void failure(Exception e, C c) {  
        handleException(e);  
    }  
});
```

# Asynchronous sockets

```
ByteBuffer buf=ByteBuffer.allocate(100);

s.read(buf, null, new CompletionHandler() {
    public void completed(Integer result, Object a) {
        buf.flip();

        r.write(buf, ...);
    }
    public void failure(Throwable t, Object a) {
        report(t);
    }
});
```

# Thread pools

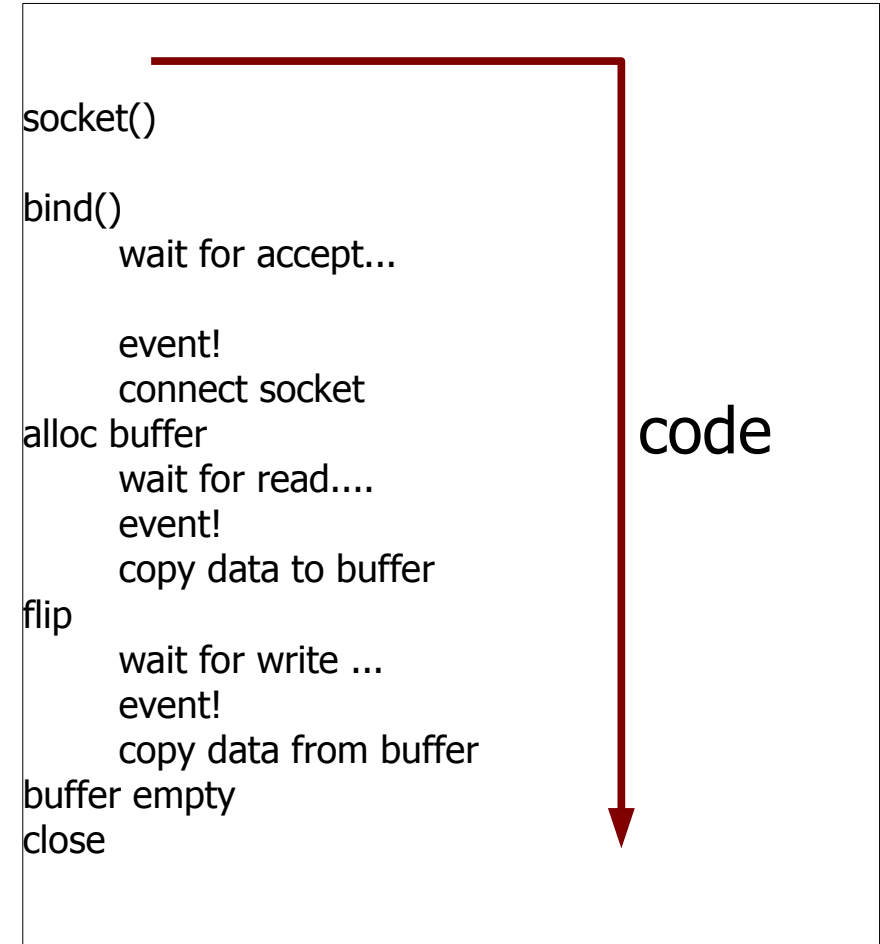
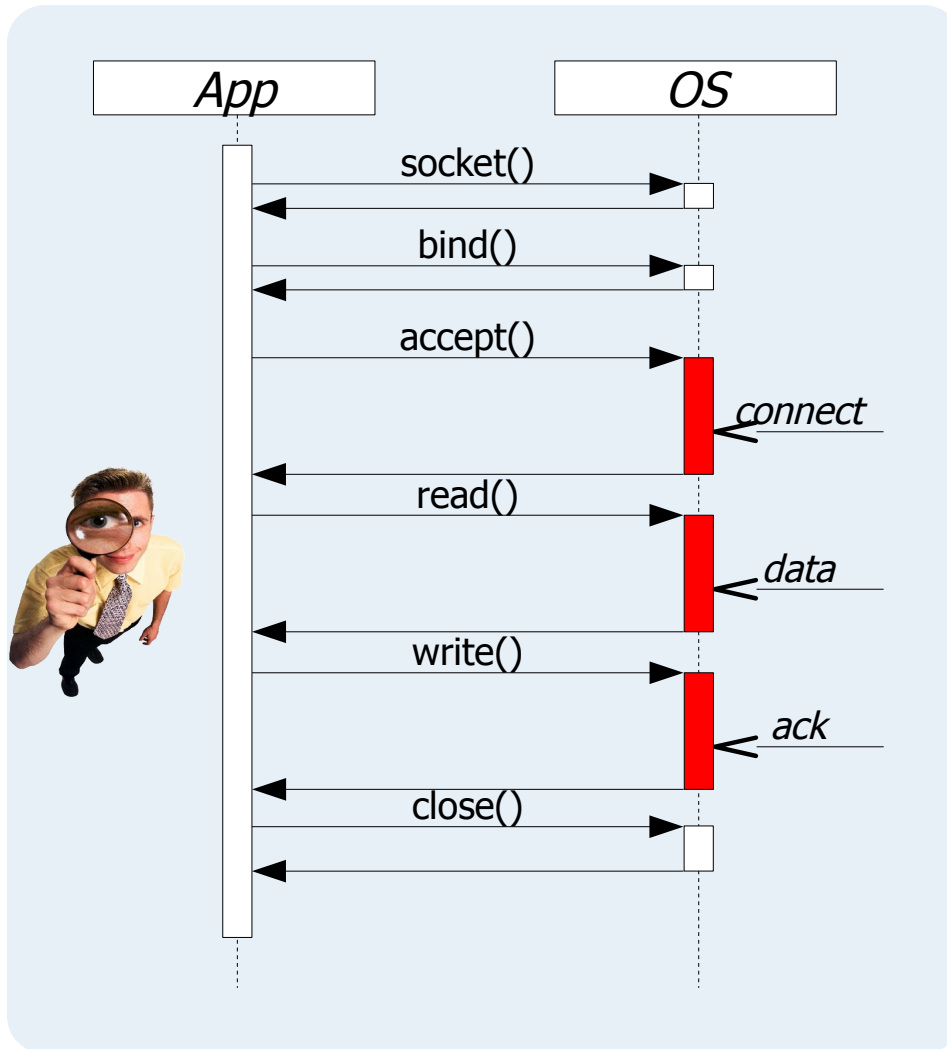
- For non-blocking, short-lived events:
  - One pool thread for hardware thread
- While all threads are blocked, the application stops handling events

```
AsynchronousChannelGroup g =  
    AsynchronousChannelGroup.withFixedThreadPool(...);  
  
AsynchronousSocketChannel s =  
    AsynchronousSocketChannel.open(g);  
  
... /* callbacks use g.shutdown() to exit */  
  
g.awaitTermination(Long.MAX_VALUE, TimeUnit.SECONDS);
```

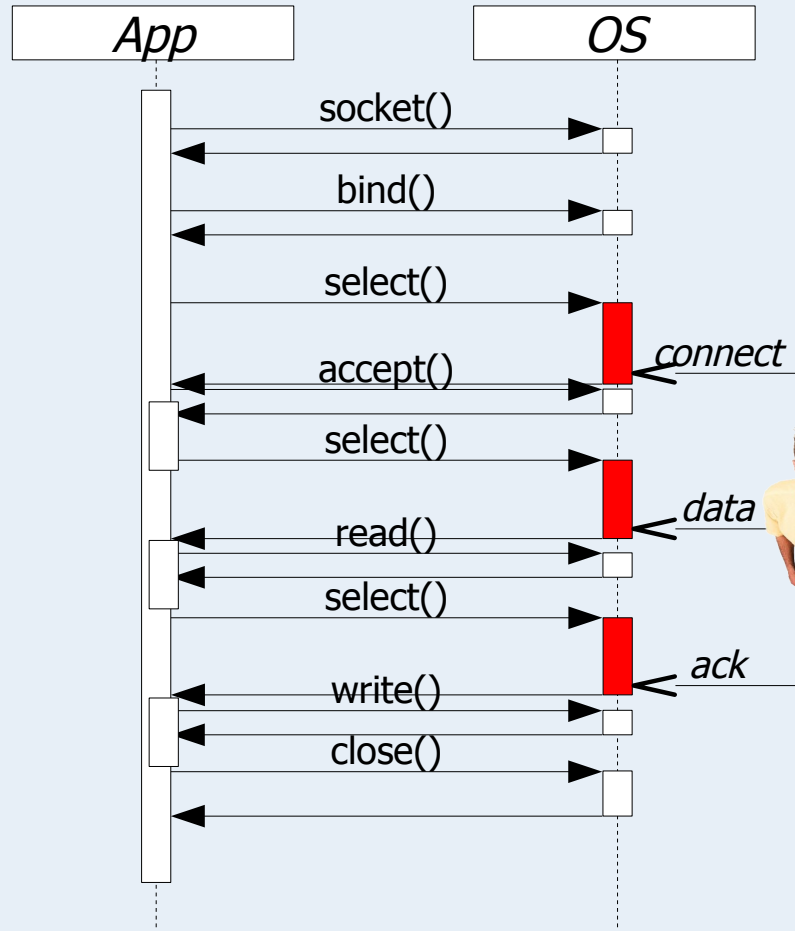
# Polled I/O

- Explicitly inform the application of which I/O channels are ready (and won't block)
  - Readiness event / level-triggered
- General idea: Instead of:  
`read(buf); doSomething();`
- do:  
`for(key: select())  
    read(buf); doSomething();`

# Threaded version



# Polled version



```

socket()
bind()
wait for accept...
event!
connect socket
alloc buffer
wait for read....
event!
alloc & copy data to buffer
flip
wait for write ...
event!
copy data from buffer
buffer empty
close
  
```

code

code

code

code



# Polled I/O

- Avoids having a dedicated thread for each event source
- Avoids captive memory for idle I/O channels
- Makes threading policy explicit
- However:
  - Requires additional system calls (and copies)
- Polled I/O in Java with NIO Selectors

# Polled I/O in Java

- Main loop:

```
Selector sel=SelectorProvider.provider().openSelector();

while(true) {
    sel.select();

    for(Iterator<SelectionKey> i=sel.selectedKeys().iterator(); i.hasNext(); ) {
        SelectionKey key = i.next();

        // i/o

        i.remove();
    }
}
```

# Polled I/O in Java

- Register interest in server socket:

```
ServerSocketChannel ss=ServerSocketChannel.open();  
ss.bind(new InetSocketAddress(12345));  
ss.configureBlocking(false);  
ss.register(sel, SelectionKey.OP_ACCEPT);
```

- Handle connection event:

```
if (key.isAcceptable()) {  
    SocketChannel s=ss.accept();  
  
    s.configureBlocking(false);  
    s.register(sel, SelectionKey.OP_READ);  
}
```

# Polled I/O in Java

```
if (key.isReadable()) {  
    ByteBuffer buf=ByteBuffer.allocate(100);  
    SocketChannel s=(SocketChannel)key.channel();  
  
    int r=s.read(buf);  
    if (r<0) {  
        key.cancel();  
        s.close();  
    } else {  
        buf.flip();  
        for(Socket r: ..) {  
            r.write(buf);  
            buf.rewind();  
        }  
    }  
}
```



What if write blocks?

# Polled I/O in Java

- Need to poll before writing
- Bytes read must be saved until writing is possible
- Signal interest on writing

```
if (key.isReadable()) {  
    ...  
} else {  
    buf.flip();  
    for(SelectionKey k: ...) {  
        key.attach(buf.duplicate());  
        key.interestOps( ... | SelectionKey.OP_WRITE);  
    }  
}  
}
```

What if multiple  
writes pending?

# Polled I/O in Java

- Get bytes attached to key
- Reset interest to reading

```
if (key.isWritable()) {  
    SocketChannel s=(SocketChannel)key.channel();  
    ByteBuffer buf=(ByteBuffer)key.attachment();  
  
    s.write(buf);  
    key.interestOps(SelectionKey.OP_READ);  
}
```

# Polled I/O + Object oriented

- Encapsulate context data + event-handling code

```
public class ChatSession implements Handler {  
    private ByteBuffer stored; // possibly a queue...  
  
    public ChatSession(SelectionKey key) {  
        // initialization  
    }  
  
    public void handleRead(ByteBuffer in) throws IOException {  
        // store input  
    }  
    public void handleWrite() throws IOException {  
        // write from stored  
    }  
}
```

# Polled I/O + Object oriented

```
if (key.isAcceptable()) {
    SocketChannel s=ss.accept();

    if (s!=null) {
        s.configureBlocking(false);
        SelectionKey nkey=s.register(sel, SelectionKey.OP_READ);
        nkey.attach(new ChatSession(...));
    }
} else if (key.isReadable()) {
    Handler h=(Handler)key.attachment();
    ByteBuffer buf = ByteBuffer.allocate(100);
    ((SocketChannel)key.channel()).read(buf);
    handler.handleRead(buf);
} else if (key.isWritable()) {
    Handler handler=(Handler)key.attachment();
    handler.handleWrite();
}
```