

Introduction to Erlang

Paulo Sérgio Almeida

Grupo de Sistemas Distribuídos
Departamento de Informática
Universidade do Minho



- Created by Ericsson; after mathematician A. K. Erlang
- Dynamically typed;
- Purely functional and strict kernel;
- Ultra-lightweight processes managed by the Erlang VM;
- Concurrency by asynchronous message passing (actor model);
- Persistent and transactional database for global state;
- emphasizes reliability; always-on systems;
- emphasizes previsibility; soft-realtime systems.



Erlang: functional languages can be used in the real world

- Even though simple and elegant, Erlang is not a toy-language;
- One of the more successful functional languages:
 - AXD 301:** Ericsson ATM switch; 1.7 Mlocs Erlang; 99.999999999% reliability (31ms/year);
 - GPRS:** 76% code in Erlang;
 - ejabberd:** Jabber/XMPP server, adopted by jabber.org; 250000 simultaneous users in a domain; 10000 users per host;
 - Yaws:** web server; worked well with 80000 connections in a test where Apache died at 4000;
 - RabbitMQ:** implementation of messaging standard AMQP (Advanced Message Queuing Protocol);
 - Riak** highly scalable, fault-tolerant distributed database;
 - WhatsApp** messaging, voice, video, . . .



- Functional language:
 - higher-order functions,
 - pattern matching,
 - no update assignment (single assignment),
 - main data structures : tuples, lists and maps;
- Dynamically typed language:
 - functional “scripting language”,
 - good for rapid prototyping,
 - static verification tools; e.g. dialyzer;



- An Erlang process is managed by the Erlang virtual machine;
- Very little creation cost;
- Easy to have hundreds of thousand processes per host;
- Very little context switch cost (one or two orders of magnitude faster than OS threads);
- Some OS threads are used by the Erlang VM to exploit parallelism or for I/O.
- Each process maintains local state and communicates by message passing with other processes.



- Erlang follows the actor model;
- Each process has a mailbox – message queue;
- Messages are sent specifying the destination process – no other concept, like channel, is used.
- Sending is asynchronous/non-blocking; it does not give an error even if the destination process has already terminated;
- Messages are kept in the mailbox by arrival order;
- Messages are removed from the mailbox by a blocking receive using pattern matching.



- Functional language with garbage collection (GC);
- Processes do not share memory;
- In a functional language with immutable data structures, absence of sharing between processes could be only conceptual . . .
- . . . but in Erlang there is really no physical sharing:
 - each process has its own stack and heap;
 - messages involve physical copy of data structures;
- Why this choice:
 - messages are typically short;
 - heap per process allows GC per process;
 - avoids large pauses caused by the GC on a global heap;
 - results in greater locality;
 - a system may be tuned so that many processes are created without incurring in a single GC during their lifetime.
- There is a shared-heap option, normally not used.



- State encapsulation in imperative languages or objects is one of the problems not yet solved;

“The big lie of object-oriented programming is that objects provide encapsulation”
John Hogg

- Mixing mutable state with functional style (e.g., in Python) brings curious surprises (bugs);
- In Erlang, a process truly encapsulates state;
- It can be seen as an object in the original sense (Alan Kay):
 - interacts by message passing;
 - encapsulates local state;



- Erlang emphasizes fast response to external events;
- Used in soft-realtime systems;
- Reflected in the language /system:
 - does not use concepts that make performance estimation difficult, like lazy evaluation;
 - ultra-lightweight processes with fast context switch;
 - per process GC to avoid long pauses;
 - in-memory database (Mnesia);



- Large systems contain bugs;
- System should work, even having some incorrect part;
- Erlang philosophy – avoid defensive programming;
 - functions do not protect against invalid arguments;
 - do not mix error verifying code with normal code;
 - program for the normal case;
- Erlang philosophy – fail-crash:
 - if function cannot perform as supposed, do not try to return error but crash process;
 - example: if guards in an “if” or patterns in a “case” do not cover all cases, process crashes;
- Processes play an essential role in error isolation and containment:
 - processes can be linked – propagating or containing failures;
 - supervisor processes can relaunch failed processes;
 - system organized as a supervision tree;



Non-stop operation and hot code swapping

- Erlang in systems that cannot be switched off or rebooted;
- Over time, new versions of some code need to be deployed;
 - bug crashes a process; error report is generated; bug is corrected;
 - new features are added to the application;
- Erlang can load new version of module in running application;
- Two ways to invoke a function:
 - intra-module; uses the same module version;
 - with module prefix; uses the more recent version;
- VM supports two versions of a module coexisting;
- Different processes can be running different versions.



- Last layer allows state that is global to all processes;
- Two modules: ets and mnesia;
- ets – Erlang Term Storage:
 - Allows tables with many entries; preferably small entries;
 - In memory, using hash tables or trees;
 - Belong to creating process; removed when it terminates;
 - May be created as private, protected, or public;
 - Kept in specific memory area, with no GC;
 - Terms are copied between processes and tables;



“Mnesia is a distributed DataBase Management System (DBMS), appropriate for telecommunications applications and other Erlang applications which require continuous operation and exhibit soft real-time properties.”

- Written in Erlang, using ets;
- Allows storing generic Erlang terms, without impedance mismatch;
- Query Language based on list comprehensions;
- Persistent; but normally an in-memory DB;
- Extremely fast reading access;
- Nested transactions using plain Erlang functions;
- Replication and fragmentation of tables;
- Control of isolation level;
- Allows reconfiguration without stopping the system.



- Current applications tend to be written using many components that must be combined;
- Impedance mismatch (e.g., object-relational) hinders development;
- Erlang + Mnesia + Yaws allows writing an application in a simple and integrated way;

“Erlang + Yaws is so much easier to setup and install than Ruby + Gems + Rails + FastCGI + Lighttpd + MySQL + MySQL bindings + every-other-package-in-the-universe. I’ve been slaving away at these installation instructions for OS X for what feels like an eternity, and my patience is running very low.”

Yariv Sadan



Dialyzer: static analysis tool to detect errors (types, unreachable code, unnecessary tests); performs data flow analysis, to determine the disjoint union of prime types;

Application Monitor: graphic tool to monitor the supervision tree of an application;

Event Tracer: to collect and visualize information using the Erlang tracing mechanism;

Process Manager: to inspect processes (local or remote) and do tracing;

Table Visualizer: to inspect or modify ets or mnesia tables;

